



In [2]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this c
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

```
/home/komalumrethe/anaconda3/lib/python3.5/site-packages/tensorflow/python/f
ramework/dtypes.py:516: FutureWarning: Passing (type, 1) or '1type' as a syn
onym of type is deprecated; in a future version of numpy, it will be underst
ood as (type, (1,)) / '(1,)type'.
```

```
_np_qint8 = np.dtype [("qint8", np.int8, 1)])
```

```
/home/komalumrethe/anaconda3/lib/python3.5/site-packages/tensorflow/python/f
ramework/dtypes.py:517: FutureWarning: Passing (type, 1) or '1type' as a syn
onym of type is deprecated; in a future version of numpy, it will be underst
ood as (type, (1,)) / '(1,)type'.
```

```
_np_quint8 = np.dtype [("quint8", np.uint8, 1)])
```

```
/home/komalumrethe/anaconda3/lib/python3.5/site-packages/tensorflow/python/f
ramework/dtypes.py:518: FutureWarning: Passing (type, 1) or '1type' as a syn
onym of type is deprecated; in a future version of numpy, it will be underst
ood as (type, (1,)) / '(1,)type'.
```

```
_np_qint16 = np.dtype [("qint16", np.int16, 1)])
```

```
/home/komalumrethe/anaconda3/lib/python3.5/site-packages/tensorflow/python/f
ramework/dtypes.py:519: FutureWarning: Passing (type, 1) or '1type' as a syn
onym of type is deprecated; in a future version of numpy, it will be underst
ood as (type, (1,)) / '(1,)type'.
```

```
_np_quint16 = np.dtype [("quint16", np.uint16, 1)])
```

```
/home/komalumrethe/anaconda3/lib/python3.5/site-packages/tensorflow/python/f
ramework/dtypes.py:520: FutureWarning: Passing (type, 1) or '1type' as a syn
onym of type is deprecated; in a future version of numpy, it will be underst
ood as (type, (1,)) / '(1,)type'.
```

```
_np_qint32 = np.dtype [("qint32", np.int32, 1)])
```

```
/home/komalumrethe/anaconda3/lib/python3.5/site-packages/tensorflow/python/f
ramework/dtypes.py:525: FutureWarning: Passing (type, 1) or '1type' as a syn
onym of type is deprecated; in a future version of numpy, it will be underst
ood as (type, (1,)) / '(1,)type'.
```

```
np_resource = np.dtype [("resource", np.ubyte, 1)])
```

```
/home/komalumrethe/anaconda3/lib/python3.5/site-packages/tensorboard/compat/
tensorflow_stub/dtypes.py:541: FutureWarning: Passing (type, 1) or '1type' a
s a synonym of type is deprecated; in a future version of numpy, it will be
understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint8 = np.dtype [("qint8", np.int8, 1)])
```

```
/home/komalumrethe/anaconda3/lib/python3.5/site-packages/tensorboard/compat/
tensorflow_stub/dtypes.py:542: FutureWarning: Passing (type, 1) or '1type' a
s a synonym of type is deprecated; in a future version of numpy, it will be
understood as (type, (1,)) / '(1,)type'.
```

```
_np_quint8 = np.dtype [("quint8", np.uint8, 1)])
```

```
/home/komalumrethe/anaconda3/lib/python3.5/site-packages/tensorboard/compat/
tensorflow_stub/dtypes.py:543: FutureWarning: Passing (type, 1) or '1type' a
s a synonym of type is deprecated; in a future version of numpy, it will be
understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint16 = np.dtype [("qint16", np.int16, 1)])
```

```
/home/komalumrethe/anaconda3/lib/python3.5/site-packages/tensorboard/compat/
tensorflow_stub/dtypes.py:544: FutureWarning: Passing (type, 1) or '1type' a
s a synonym of type is deprecated; in a future version of numpy, it will be
understood as (type, (1,)) / '(1,)type'.
```

```
_np_quint16 = np.dtype [("quint16", np.uint16, 1)])
```

```
/home/komalumrethe/anaconda3/lib/python3.5/site-packages/tensorboard/compat/
tensorflow_stub/dtypes.py:545: FutureWarning: Passing (type, 1) or '1type' a
```

```
s a synonym of type is deprecated; in a future version of numpy, it will be
understood as (type, (1,)) / '(1,)type'.
_np_qint32 = np.dtype(["qint32", np.int32, 1])
/home/komalumrethe/anaconda3/lib/python3.5/site-packages/tensorboard/compat/
tensorflow_stub/dtypes.py:550: FutureWarning: Passing (type, 1) or '1type' a
s a synonym of type is deprecated; in a future version of numpy, it will be
understood as (type, (1,)) / '(1,)type'.
_np_resource = np.dtype(["resource", np.ubyte, 1])
```

In [3]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    # plt.show()
    # fig.canvas.draw()
    plt.show()
```

In [4]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz> (<http://s3.amazonaws.com/img-datasets/mnist.npz>)  
11493376/11490434 [=====] - 0s 0us/step

In [5]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)" % (X_train.shape[1], X_train.shape[2]))
print("Number of testing examples :", X_test.shape[0], "and each image is of shape (%d, %d)" % (X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)  
Number of testing examples : 10000 and each image is of shape (28, 28)

In [6]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [7]:

```
# after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)%"
print("Number of testing examples :", X_test.shape[0], "and each image is of shape (%d)%"(X
```

```
Number of training examples : 60000 and each image is of shape (784)
```

```
Number of testing examples : 10000 and each image is of shape (784)
```

In [8]:

```
# An example data point
```

```
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18  126  136  175  26  166  255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94  154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251 93  82
 82 56 39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205 11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253 90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  11 190 253 70  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119 25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150 27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  16 93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  249 253 249 64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201 78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0 23 66 213 253 253 253 253 198 81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 195
80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
55 172 226 253 253 253 253 244 133 11  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  136 253 253 253 212 135 132 16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
```

In [9]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
X_test = X_test/255
```

In [10]:

```
# example data point after normlizing
print(X_train[0])
```

```
[0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.]
```

In [11]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# Lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

Class label of first image : 5

After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

## Softmax classifier

In [12]:

```
# https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT.X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument of a layer.

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation
```

In [13]:

```
# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

## **MLP + ReLU + Adam: 2 Hidden Layers and without Dropout and Batch Normalization**

In [14]:

```

model_relu = Sequential()
model_relu.add(Dense(364, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_relu.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)

```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 364)	285740
dense_2 (Dense)	(None, 52)	18980
dense_3 (Dense)	(None, 10)	530
Total params: 305,250		
Trainable params: 305,250		
Non-trainable params: 0		

None

WARNING:tensorflow:From /home/komalumrethe/anaconda3/lib/python3.5/site-packages/keras/backend/tensorflow\_backend.py:422: The name tf.global\_variables is deprecated. Please use tf.compat.v1.global\_variables instead.

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 39us/step - loss: 0.2553  
 - accuracy: 0.9239 - val\_loss: 0.1404 - val\_accuracy: 0.9576

Epoch 2/20

60000/60000 [=====] - 2s 36us/step - loss: 0.0980  
 - accuracy: 0.9707 - val\_loss: 0.1119 - val\_accuracy: 0.9661

Epoch 3/20

60000/60000 [=====] - 2s 35us/step - loss: 0.0620  
 - accuracy: 0.9810 - val\_loss: 0.0757 - val\_accuracy: 0.9770

Epoch 4/20

60000/60000 [=====] - 2s 35us/step - loss: 0.0450  
 - accuracy: 0.9859 - val\_loss: 0.0855 - val\_accuracy: 0.9742

Epoch 5/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0312  
 - accuracy: 0.9907 - val\_loss: 0.0774 - val\_accuracy: 0.9784

Epoch 6/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0240  
 - accuracy: 0.9926 - val\_loss: 0.0749 - val\_accuracy: 0.9773

Epoch 7/20

60000/60000 [=====] - 2s 36us/step - loss: 0.0185  
 - accuracy: 0.9942 - val\_loss: 0.0788 - val\_accuracy: 0.9789

Epoch 8/20

60000/60000 [=====] - 2s 35us/step - loss: 0.0139  
 - accuracy: 0.9961 - val\_loss: 0.0794 - val\_accuracy: 0.9783

Epoch 9/20

60000/60000 [=====] - 2s 36us/step - loss: 0.0153  
 - accuracy: 0.9952 - val\_loss: 0.0907 - val\_accuracy: 0.9759



```
Epoch 10/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0112
- accuracy: 0.9968 - val_loss: 0.0777 - val_accuracy: 0.9795
Epoch 11/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0083
- accuracy: 0.9976 - val_loss: 0.0845 - val_accuracy: 0.9799
Epoch 12/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0126
- accuracy: 0.9960 - val_loss: 0.0876 - val_accuracy: 0.9779
Epoch 13/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0103
- accuracy: 0.9966 - val_loss: 0.0890 - val_accuracy: 0.9786
Epoch 14/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0078
- accuracy: 0.9976 - val_loss: 0.0864 - val_accuracy: 0.9797
Epoch 15/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0054
- accuracy: 0.9983 - val_loss: 0.0876 - val_accuracy: 0.9817
Epoch 16/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0129
- accuracy: 0.9957 - val_loss: 0.0881 - val_accuracy: 0.9813
Epoch 17/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0083
- accuracy: 0.9973 - val_loss: 0.0890 - val_accuracy: 0.9800
Epoch 18/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0046
- accuracy: 0.9985 - val_loss: 0.0856 - val_accuracy: 0.9815
Epoch 19/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0053
- accuracy: 0.9982 - val_loss: 0.0905 - val_accuracy: 0.9815
Epoch 20/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0106
- accuracy: 0.9968 - val_loss: 0.0957 - val_accuracy: 0.9792
```

In [15]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

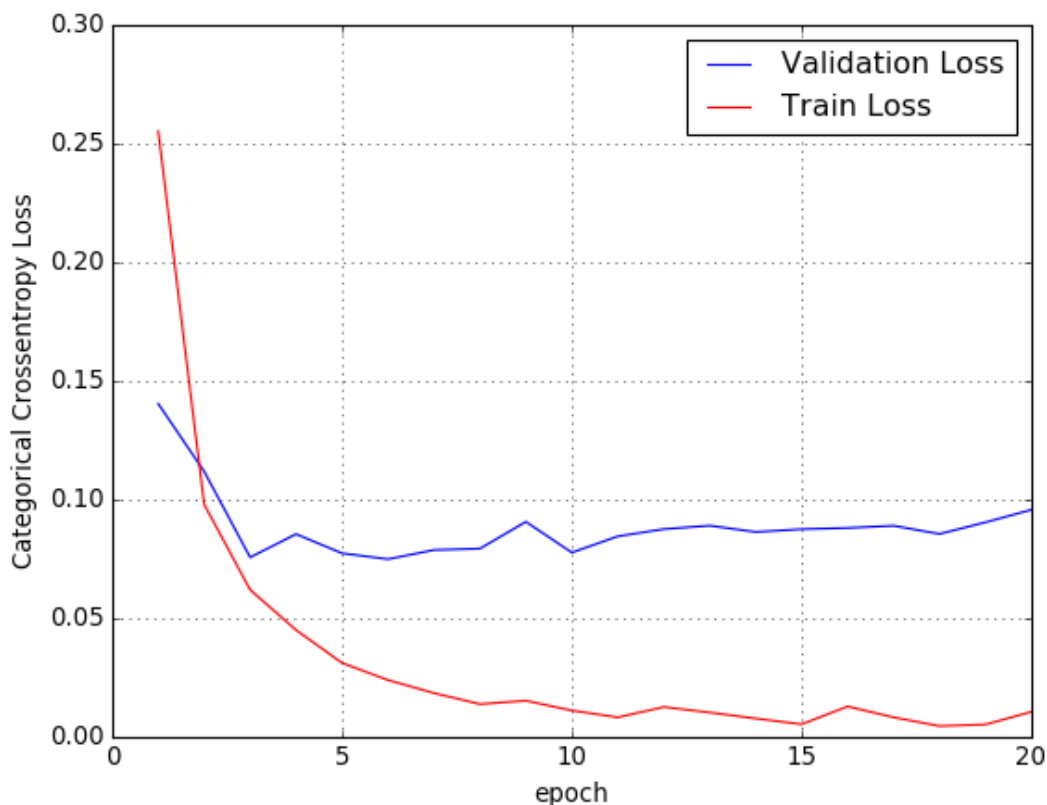
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09574766628526103

Test accuracy: 0.979200005531311



In [16]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## MLP + Batch-Norm on 2 hidden Layers + AdamOptimizer

In [17]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(0, \sigma) = N(0, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(0, \sigma) = N(0, 0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(364, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.120, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
dense_4 (Dense)	(None, 364)	285740
batch_normalization_1 (Batch Normalization)	(None, 364)	1456
dense_5 (Dense)	(None, 52)	18980
batch_normalization_2 (Batch Normalization)	(None, 52)	208
dense_6 (Dense)	(None, 10)	530
=====		
Total params: 306,914		
Trainable params: 306,082		
Non-trainable params: 832		

In [18]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                          validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 62us/step - loss: 0.2564 - accuracy: 0.9274 - val\_loss: 0.1154 - val\_accuracy: 0.9668

Epoch 2/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0888 - accuracy: 0.9743 - val\_loss: 0.0926 - val\_accuracy: 0.9719

Epoch 3/20

60000/60000 [=====] - 3s 48us/step - loss: 0.0560 - accuracy: 0.9835 - val\_loss: 0.0806 - val\_accuracy: 0.9741

Epoch 4/20

60000/60000 [=====] - 3s 47us/step - loss: 0.0393 - accuracy: 0.9876 - val\_loss: 0.0829 - val\_accuracy: 0.9733

Epoch 5/20

60000/60000 [=====] - 3s 47us/step - loss: 0.0269 - accuracy: 0.9919 - val\_loss: 0.0819 - val\_accuracy: 0.9741

Epoch 6/20

60000/60000 [=====] - 3s 48us/step - loss: 0.0212 - accuracy: 0.9935 - val\_loss: 0.0815 - val\_accuracy: 0.9748

Epoch 7/20

60000/60000 [=====] - 3s 48us/step - loss: 0.0173 - accuracy: 0.9950 - val\_loss: 0.0905 - val\_accuracy: 0.9736

Epoch 8/20

60000/60000 [=====] - 3s 48us/step - loss: 0.0154 - accuracy: 0.9952 - val\_loss: 0.0847 - val\_accuracy: 0.9769

Epoch 9/20

60000/60000 [=====] - 3s 48us/step - loss: 0.0137 - accuracy: 0.9958 - val\_loss: 0.0884 - val\_accuracy: 0.9758

Epoch 10/20

60000/60000 [=====] - 3s 47us/step - loss: 0.0141 - accuracy: 0.9955 - val\_loss: 0.0831 - val\_accuracy: 0.9762

Epoch 11/20

60000/60000 [=====] - 3s 48us/step - loss: 0.0106 - accuracy: 0.9969 - val\_loss: 0.0867 - val\_accuracy: 0.9764

Epoch 12/20

60000/60000 [=====] - 3s 48us/step - loss: 0.0131 - accuracy: 0.9956 - val\_loss: 0.0825 - val\_accuracy: 0.9771

Epoch 13/20

60000/60000 [=====] - 3s 47us/step - loss: 0.0088 - accuracy: 0.9973 - val\_loss: 0.0877 - val\_accuracy: 0.9769

Epoch 14/20

60000/60000 [=====] - 3s 47us/step - loss: 0.0097 - accuracy: 0.9969 - val\_loss: 0.0880 - val\_accuracy: 0.9785

Epoch 15/20

60000/60000 [=====] - 3s 47us/step - loss: 0.0084 - accuracy: 0.9974 - val\_loss: 0.0905 - val\_accuracy: 0.9761

Epoch 16/20

60000/60000 [=====] - 3s 47us/step - loss: 0.0080 - accuracy: 0.9975 - val\_loss: 0.0847 - val\_accuracy: 0.9776

Epoch 17/20

60000/60000 [=====] - 3s 47us/step - loss: 0.0071 - accuracy: 0.9976 - val\_loss: 0.0880 - val\_accuracy: 0.9786

Epoch 18/20

60000/60000 [=====] - 3s 47us/step - loss: 0.0079 -

accuracy: 0.9975 - val\_loss: 0.0839 - val\_accuracy: 0.9805

Epoch 19/20

60000/60000 [=====] - 3s 46us/step - loss: 0.0063 -

accuracy: 0.9980 - val\_loss: 0.0828 - val\_accuracy: 0.9792

Epoch 20/20

60000/60000 [=====] - 3s 47us/step - loss: 0.0063 -

accuracy: 0.9980 - val\_loss: 0.0872 - val\_accuracy: 0.9796

In [19]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0872070386303676

Test accuracy: 0.9796000123023987



In [20]:

```

w_after = model_batch.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



## MLP + Dropout (dropout rate = 0.5) on 2 hidden layers +



# AdamOptimizer

In [21]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-funct

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
=====		
dense_7 (Dense)	(None, 364)	285740
-----		
batch_normalization_3 (Batch Normalization)	(None, 364)	1456
-----		
dropout_1 (Dropout)	(None, 364)	0
-----		
dense_8 (Dense)	(None, 52)	18980
-----		
batch_normalization_4 (Batch Normalization)	(None, 52)	208
-----		
dropout_2 (Dropout)	(None, 52)	0
-----		
dense_9 (Dense)	(None, 10)	530
=====		
Total params: 306,914		
Trainable params: 306,082		
Non-trainable params: 832		

In [22]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
                        validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 63us/step - loss: 0.6613 - accuracy: 0.7983 - val\_loss: 0.1981 - val\_accuracy: 0.9418

Epoch 2/20

60000/60000 [=====] - 3s 58us/step - loss: 0.3129 - accuracy: 0.9081 - val\_loss: 0.1463 - val\_accuracy: 0.9566

Epoch 3/20

60000/60000 [=====] - 3s 56us/step - loss: 0.2397 - accuracy: 0.9315 - val\_loss: 0.1147 - val\_accuracy: 0.9666

Epoch 4/20

60000/60000 [=====] - 3s 54us/step - loss: 0.2024 - accuracy: 0.9421 - val\_loss: 0.1061 - val\_accuracy: 0.9666

Epoch 5/20

60000/60000 [=====] - 3s 57us/step - loss: 0.1766 - accuracy: 0.9488 - val\_loss: 0.0959 - val\_accuracy: 0.9709

Epoch 6/20

60000/60000 [=====] - 3s 55us/step - loss: 0.1585 - accuracy: 0.9534 - val\_loss: 0.0912 - val\_accuracy: 0.9731

Epoch 7/20

60000/60000 [=====] - 3s 54us/step - loss: 0.1439 - accuracy: 0.9577 - val\_loss: 0.0841 - val\_accuracy: 0.9746

Epoch 8/20

60000/60000 [=====] - 3s 53us/step - loss: 0.1366 - accuracy: 0.9604 - val\_loss: 0.0823 - val\_accuracy: 0.9750

Epoch 9/20

60000/60000 [=====] - 3s 52us/step - loss: 0.1233 - accuracy: 0.9642 - val\_loss: 0.0749 - val\_accuracy: 0.9764

Epoch 10/20

60000/60000 [=====] - 3s 53us/step - loss: 0.1192 - accuracy: 0.9650 - val\_loss: 0.0701 - val\_accuracy: 0.9790

Epoch 11/20

60000/60000 [=====] - 3s 54us/step - loss: 0.1116 - accuracy: 0.9664 - val\_loss: 0.0718 - val\_accuracy: 0.9789

Epoch 12/20

60000/60000 [=====] - 3s 54us/step - loss: 0.1080 - accuracy: 0.9690 - val\_loss: 0.0715 - val\_accuracy: 0.9788

Epoch 13/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0971 - accuracy: 0.9713 - val\_loss: 0.0685 - val\_accuracy: 0.9796

Epoch 14/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0969 - accuracy: 0.9706 - val\_loss: 0.0689 - val\_accuracy: 0.9804

Epoch 15/20

60000/60000 [=====] - 4s 60us/step - loss: 0.0928 - accuracy: 0.9726 - val\_loss: 0.0695 - val\_accuracy: 0.9801

Epoch 16/20

60000/60000 [=====] - 4s 62us/step - loss: 0.0880 - accuracy: 0.9736 - val\_loss: 0.0690 - val\_accuracy: 0.9799

Epoch 17/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0857 - accuracy: 0.9744 - val\_loss: 0.0652 - val\_accuracy: 0.9806

Epoch 18/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0819 -

```

accuracy: 0.9756 - val_loss: 0.0671 - val_accuracy: 0.9818
Epoch 19/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0808 -
accuracy: 0.9759 - val_loss: 0.0657 - val_accuracy: 0.9803
Epoch 20/20
60000/60000 [=====] - 3s 54us/step - loss: 0.0751 -
accuracy: 0.9775 - val_loss: 0.0625 - val_accuracy: 0.9828

```

In [ ]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

## MLP + Dropout (dropout rate = 0.25) on 2 hidden layers + AdamOptimizer

In [24]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-funct
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.25))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.25))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
=====		
dense_13 (Dense)	(None, 364)	285740
batch_normalization_7 (Batch Normalization)	(None, 364)	1456
dropout_5 (Dropout)	(None, 364)	0
dense_14 (Dense)	(None, 52)	18980
batch_normalization_8 (Batch Normalization)	(None, 52)	208
dropout_6 (Dropout)	(None, 52)	0
dense_15 (Dense)	(None, 10)	530
=====		
Total params: 306,914		
Trainable params: 306,082		
Non-trainable params: 832		

In [25]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
                        validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 64us/step - loss: 0.3875 - accuracy: 0.8858 - val\_loss: 0.1491 - val\_accuracy: 0.9571

Epoch 2/20

60000/60000 [=====] - 3s 54us/step - loss: 0.1723 - accuracy: 0.9498 - val\_loss: 0.1036 - val\_accuracy: 0.9681

Epoch 3/20

60000/60000 [=====] - 3s 53us/step - loss: 0.1295 - accuracy: 0.9614 - val\_loss: 0.0851 - val\_accuracy: 0.9737

Epoch 4/20

60000/60000 [=====] - 3s 53us/step - loss: 0.1073 - accuracy: 0.9673 - val\_loss: 0.0894 - val\_accuracy: 0.9728

Epoch 5/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0900 - accuracy: 0.9717 - val\_loss: 0.0730 - val\_accuracy: 0.9784

Epoch 6/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0771 - accuracy: 0.9762 - val\_loss: 0.0706 - val\_accuracy: 0.9771

Epoch 7/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0724 - accuracy: 0.9771 - val\_loss: 0.0718 - val\_accuracy: 0.9779

Epoch 8/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0652 - accuracy: 0.9793 - val\_loss: 0.0682 - val\_accuracy: 0.9783

Epoch 9/20

60000/60000 [=====] - 4s 60us/step - loss: 0.0565 - accuracy: 0.9819 - val\_loss: 0.0669 - val\_accuracy: 0.9807

Epoch 10/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0546 - accuracy: 0.9827 - val\_loss: 0.0658 - val\_accuracy: 0.9805

Epoch 11/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0483 - accuracy: 0.9842 - val\_loss: 0.0643 - val\_accuracy: 0.9810

Epoch 12/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0461 - accuracy: 0.9854 - val\_loss: 0.0673 - val\_accuracy: 0.9815

Epoch 13/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0418 - accuracy: 0.9862 - val\_loss: 0.0693 - val\_accuracy: 0.9805

Epoch 14/20

60000/60000 [=====] - 4s 59us/step - loss: 0.0403 - accuracy: 0.9868 - val\_loss: 0.0652 - val\_accuracy: 0.9809

Epoch 15/20

60000/60000 [=====] - 4s 65us/step - loss: 0.0389 - accuracy: 0.9872 - val\_loss: 0.0608 - val\_accuracy: 0.9817

Epoch 16/20

60000/60000 [=====] - 4s 65us/step - loss: 0.0373 - accuracy: 0.9876 - val\_loss: 0.0585 - val\_accuracy: 0.9836

Epoch 17/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0368 - accuracy: 0.9876 - val\_loss: 0.0629 - val\_accuracy: 0.9832

Epoch 18/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0325 -

accuracy: 0.9892 - val\_loss: 0.0611 - val\_accuracy: 0.9824

Epoch 19/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0294 -

accuracy: 0.9902 - val\_loss: 0.0637 - val\_accuracy: 0.9826

Epoch 20/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0292 -

accuracy: 0.9907 - val\_loss: 0.0636 - val\_accuracy: 0.9829

## MLP + Dropout (dropout rate = 0.75) on 2 hidden layers + AdamOptimizer

In [27]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-funct
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.25))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.25))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Model: "sequential\_6"

Layer (type)	Output Shape	Param #
dense_16 (Dense)	(None, 364)	285740
batch_normalization_9 (Batch Normalization)	(None, 364)	1456
dropout_7 (Dropout)	(None, 364)	0
dense_17 (Dense)	(None, 52)	18980
batch_normalization_10 (Batch Normalization)	(None, 52)	208
dropout_8 (Dropout)	(None, 52)	0
dense_18 (Dense)	(None, 10)	530
Total params: 306,914		
Trainable params: 306,082		
Non-trainable params: 832		

In [28]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
                        validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 67us/step - loss: 0.4016 - accuracy: 0.8825 - val\_loss: 0.1430 - val\_accuracy: 0.9583

Epoch 2/20

60000/60000 [=====] - 3s 55us/step - loss: 0.1653 - accuracy: 0.9513 - val\_loss: 0.1010 - val\_accuracy: 0.9689

Epoch 3/20

60000/60000 [=====] - 3s 54us/step - loss: 0.1259 - accuracy: 0.9630 - val\_loss: 0.0912 - val\_accuracy: 0.9723

Epoch 4/20

60000/60000 [=====] - 3s 55us/step - loss: 0.1010 - accuracy: 0.9685 - val\_loss: 0.0823 - val\_accuracy: 0.9741

Epoch 5/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0858 - accuracy: 0.9734 - val\_loss: 0.0733 - val\_accuracy: 0.9782

Epoch 6/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0745 - accuracy: 0.9769 - val\_loss: 0.0722 - val\_accuracy: 0.9794

Epoch 7/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0656 - accuracy: 0.9792 - val\_loss: 0.0679 - val\_accuracy: 0.9793

Epoch 8/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0598 - accuracy: 0.9807 - val\_loss: 0.0650 - val\_accuracy: 0.9813

Epoch 9/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0567 - accuracy: 0.9817 - val\_loss: 0.0698 - val\_accuracy: 0.9795

Epoch 10/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0508 - accuracy: 0.9839 - val\_loss: 0.0667 - val\_accuracy: 0.9813

Epoch 11/20

60000/60000 [=====] - 4s 62us/step - loss: 0.0452 - accuracy: 0.9857 - val\_loss: 0.0627 - val\_accuracy: 0.9822

Epoch 12/20

60000/60000 [=====] - 4s 67us/step - loss: 0.0434 - accuracy: 0.9857 - val\_loss: 0.0634 - val\_accuracy: 0.9821

Epoch 13/20

60000/60000 [=====] - 4s 64us/step - loss: 0.0402 - accuracy: 0.9870 - val\_loss: 0.0605 - val\_accuracy: 0.9815

Epoch 14/20

60000/60000 [=====] - 3s 58us/step - loss: 0.0385 - accuracy: 0.9869 - val\_loss: 0.0638 - val\_accuracy: 0.9822

Epoch 15/20

60000/60000 [=====] - 4s 59us/step - loss: 0.0352 - accuracy: 0.9883 - val\_loss: 0.0644 - val\_accuracy: 0.9825

Epoch 16/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0334 - accuracy: 0.9892 - val\_loss: 0.0665 - val\_accuracy: 0.9812

Epoch 17/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0311 - accuracy: 0.9899 - val\_loss: 0.0667 - val\_accuracy: 0.9829

Epoch 18/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0301 -

accuracy: 0.9901 - val\_loss: 0.0639 - val\_accuracy: 0.9829

Epoch 19/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0301 -

accuracy: 0.9899 - val\_loss: 0.0676 - val\_accuracy: 0.9821

Epoch 20/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0299 -

accuracy: 0.9898 - val\_loss: 0.0614 - val\_accuracy: 0.9835



In [29]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.061442252652550815  
 Test accuracy: 0.9835000038146973



In [30]:


```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## MLP + ReLU + ADAM with 3 hidden layers without Dropout and

# Batch Normalisation

In [31]:

```
model_relu = Sequential()
model_relu.add(Dense(364, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_relu.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                        validation_data=(X_test, Y_test))
```

Model: "sequential\_7"

Layer (type)	Output Shape	Param #
dense_19 (Dense)	(None, 364)	285740
dense_20 (Dense)	(None, 128)	46720
dense_21 (Dense)	(None, 52)	6708
dense_22 (Dense)	(None, 10)	530
Total params: 339,698		
Trainable params: 339,698		
Non-trainable params: 0		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 59us/step - loss: 0.2693 - accuracy: 0.9198 - val\_loss: 0.1260 - val\_accuracy: 0.9616

Epoch 2/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0999 - accuracy: 0.9691 - val\_loss: 0.1039 - val\_accuracy: 0.9661

Epoch 3/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0611 - accuracy: 0.9812 - val\_loss: 0.0795 - val\_accuracy: 0.9751

Epoch 4/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0442 - accuracy: 0.9858 - val\_loss: 0.0860 - val\_accuracy: 0.9731

Epoch 5/20

60000/60000 [=====] - 3s 58us/step - loss: 0.0307 - accuracy: 0.9900 - val\_loss: 0.0898 - val\_accuracy: 0.9751

Epoch 6/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0277 - accuracy: 0.9909 - val\_loss: 0.1259 - val\_accuracy: 0.9657

Epoch 7/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0227 - accuracy: 0.9923 - val\_loss: 0.0904 - val\_accuracy: 0.9763

Epoch 8/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0165 - accuracy: 0.9947 - val\_loss: 0.0880 - val\_accuracy: 0.9757

Epoch 9/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0153 - accuracy: 0.9951 - val\_loss: 0.0949 - val\_accuracy: 0.9758

Epoch 10/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0172 - accuracy: 0.9946 - val\_loss: 0.0866 - val\_accuracy: 0.9773

Epoch 11/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0142 - accuracy: 0.9952 - val\_loss: 0.0991 - val\_accuracy: 0.9766

Epoch 12/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0131 - accuracy: 0.9956 - val\_loss: 0.1174 - val\_accuracy: 0.9758

Epoch 13/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0163 - accuracy: 0.9949 - val\_loss: 0.1040 - val\_accuracy: 0.9783

Epoch 14/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0097 - accuracy: 0.9965 - val\_loss: 0.1023 - val\_accuracy: 0.9775

Epoch 15/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0121 - accuracy: 0.9962 - val\_loss: 0.1054 - val\_accuracy: 0.9767

Epoch 16/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0110 - accuracy: 0.9963 - val\_loss: 0.1025 - val\_accuracy: 0.9787

Epoch 17/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0080 - accuracy: 0.9973 - val\_loss: 0.1088 - val\_accuracy: 0.9776

Epoch 18/20

60000/60000 [=====] - 4s 59us/step - loss: 0.0087 - accuracy: 0.9973 - val\_loss: 0.1118 - val\_accuracy: 0.9769

Epoch 19/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0106 - accuracy: 0.9964 - val\_loss: 0.1027 - val\_accuracy: 0.9794

Epoch 20/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0067 - accuracy: 0.9976 - val\_loss: 0.1126 - val\_accuracy: 0.9791

In [32]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

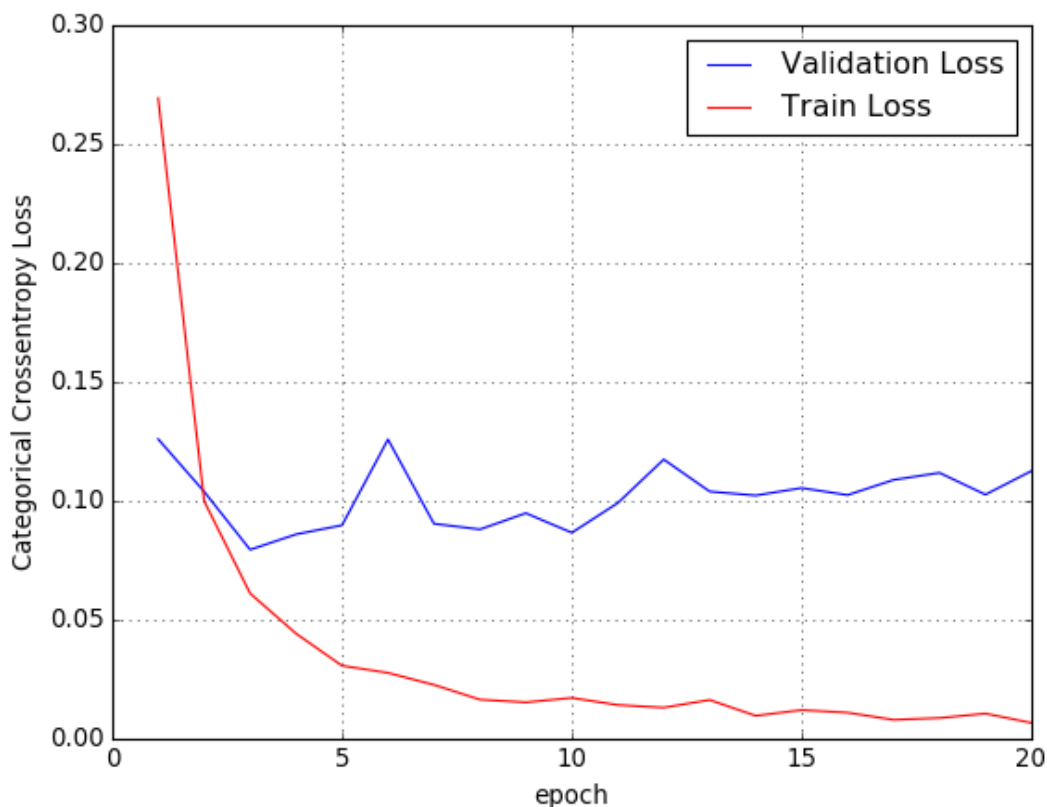
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.11258738235081656

Test accuracy: 0.9790999889373779



In [35]:

```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

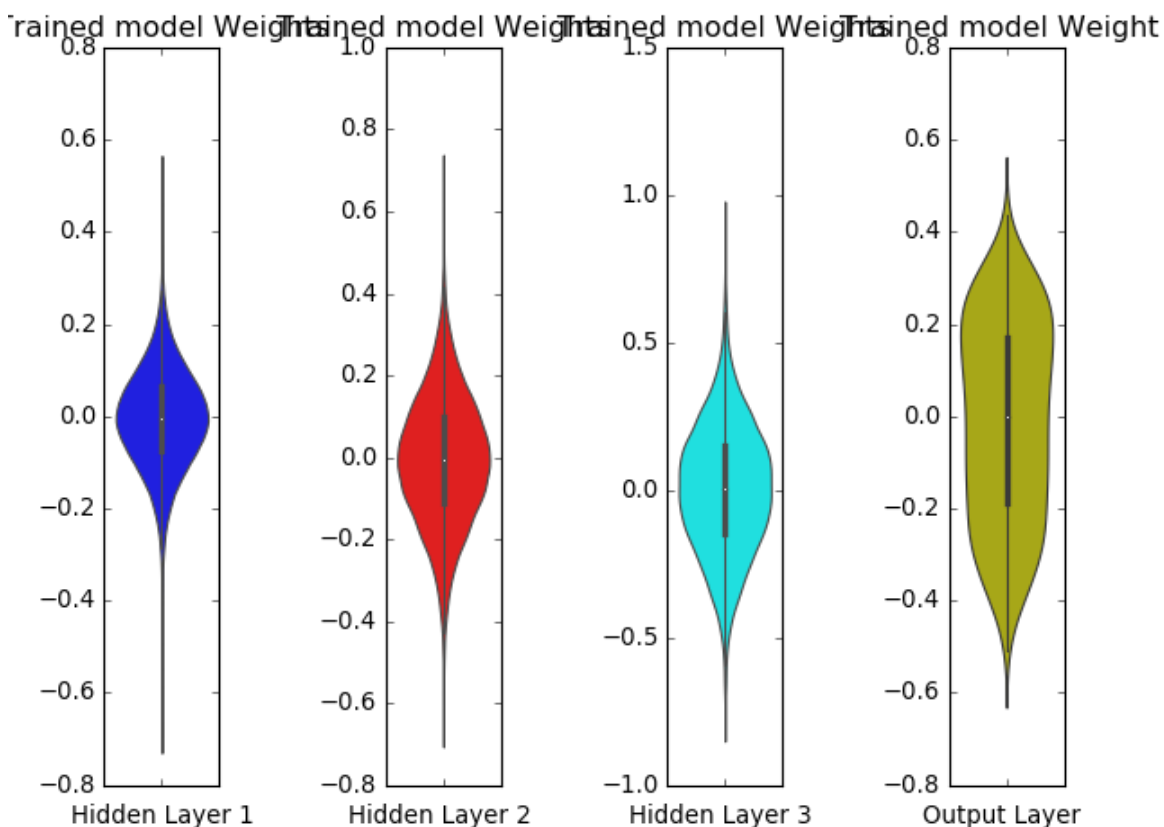
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='cyan')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
plt.tight_layout()

```



## MLP + ReLU + ADAM with 3 hidden layers with Batch Normalisation

In [36]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(0, \sigma) = N(0, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(0, \sigma) = N(0, 0.055)$ 
# h3 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(364, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.120, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.055, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Model: "sequential\_8"

Layer (type)	Output Shape	Param #
=====		
dense_23 (Dense)	(None, 364)	285740
batch_normalization_11 (Batch Normalization)	(None, 364)	1456
dense_24 (Dense)	(None, 128)	46720
batch_normalization_12 (Batch Normalization)	(None, 128)	512
dense_25 (Dense)	(None, 52)	6708
batch_normalization_13 (Batch Normalization)	(None, 52)	208
dense_26 (Dense)	(None, 10)	530
=====		
Total params: 341,874		
Trainable params: 340,786		
Non-trainable params: 1,088		



In [37]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                          validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 97us/step - loss: 0.2511  
- accuracy: 0.9278 - val\_loss: 0.1079 - val\_accuracy: 0.9677

Epoch 2/20

60000/60000 [=====] - 5s 82us/step - loss: 0.0846  
- accuracy: 0.9751 - val\_loss: 0.0883 - val\_accuracy: 0.9731

Epoch 3/20

60000/60000 [=====] - 5s 83us/step - loss: 0.0529  
- accuracy: 0.9836 - val\_loss: 0.0844 - val\_accuracy: 0.9727

Epoch 4/20

60000/60000 [=====] - 5s 80us/step - loss: 0.0363  
- accuracy: 0.9890 - val\_loss: 0.0739 - val\_accuracy: 0.9762

Epoch 5/20

60000/60000 [=====] - 5s 79us/step - loss: 0.0277  
- accuracy: 0.9913 - val\_loss: 0.0801 - val\_accuracy: 0.9765

Epoch 6/20

60000/60000 [=====] - 5s 83us/step - loss: 0.0218  
- accuracy: 0.9936 - val\_loss: 0.0795 - val\_accuracy: 0.9765

Epoch 7/20

60000/60000 [=====] - 5s 80us/step - loss: 0.0227  
- accuracy: 0.9924 - val\_loss: 0.0825 - val\_accuracy: 0.9775

Epoch 8/20

60000/60000 [=====] - 5s 83us/step - loss: 0.0169  
- accuracy: 0.9944 - val\_loss: 0.0696 - val\_accuracy: 0.9796

Epoch 9/20

60000/60000 [=====] - 5s 80us/step - loss: 0.0157  
- accuracy: 0.9949 - val\_loss: 0.0819 - val\_accuracy: 0.9760

Epoch 10/20

60000/60000 [=====] - 5s 81us/step - loss: 0.0155  
- accuracy: 0.9946 - val\_loss: 0.0790 - val\_accuracy: 0.9780

Epoch 11/20

60000/60000 [=====] - 5s 81us/step - loss: 0.0114  
- accuracy: 0.9962 - val\_loss: 0.0781 - val\_accuracy: 0.9788

Epoch 12/20

60000/60000 [=====] - 5s 80us/step - loss: 0.0111  
- accuracy: 0.9963 - val\_loss: 0.0817 - val\_accuracy: 0.9788

Epoch 13/20

60000/60000 [=====] - 5s 82us/step - loss: 0.0110  
- accuracy: 0.9964 - val\_loss: 0.0765 - val\_accuracy: 0.9801

Epoch 14/20

60000/60000 [=====] - 5s 82us/step - loss: 0.0101  
- accuracy: 0.9969 - val\_loss: 0.0767 - val\_accuracy: 0.9799

Epoch 15/20

60000/60000 [=====] - 5s 83us/step - loss: 0.0100  
- accuracy: 0.9967 - val\_loss: 0.0805 - val\_accuracy: 0.9788

Epoch 16/20

60000/60000 [=====] - 5s 81us/step - loss: 0.0092  
- accuracy: 0.9966 - val\_loss: 0.0817 - val\_accuracy: 0.9794

Epoch 17/20

60000/60000 [=====] - 5s 81us/step - loss: 0.0101  
- accuracy: 0.9967 - val\_loss: 0.0915 - val\_accuracy: 0.9770

Epoch 18/20

60000/60000 [=====] - 5s 80us/step - loss: 0.0108

```
- accuracy: 0.9962 - val_loss: 0.0818 - val_accuracy: 0.9786
Epoch 19/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0074
- accuracy: 0.9977 - val_loss: 0.0864 - val_accuracy: 0.9802
Epoch 20/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0066
- accuracy: 0.9979 - val_loss: 0.0847 - val_accuracy: 0.9790
```

In [38]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08473182359161073  
 Test accuracy: 0.9789999723434448



In [39]:

```

w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



## **MLP + ReLU + ADAM with 3 hidden layers with Dropout (dropout rate = 0.5)**

In [40]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-funct
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Model: "sequential\_9"

Layer (type)	Output Shape	Param #
dense_27 (Dense)	(None, 364)	285740
batch_normalization_14 (Batch Normalization)	(None, 364)	1456
dropout_9 (Dropout)	(None, 364)	0
dense_28 (Dense)	(None, 128)	46720
batch_normalization_15 (Batch Normalization)	(None, 128)	512
dropout_10 (Dropout)	(None, 128)	0
dense_29 (Dense)	(None, 52)	6708
batch_normalization_16 (Batch Normalization)	(None, 52)	208
dropout_11 (Dropout)	(None, 52)	0
dense_30 (Dense)	(None, 10)	530
Total params: 341,874		
Trainable params: 340,786		
Non-trainable params: 1,088		

In [41]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
                        validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 110us/step - loss: 0.9108 - accuracy: 0.7163 - val\_loss: 0.2433 - val\_accuracy: 0.9255

Epoch 2/20

60000/60000 [=====] - 6s 92us/step - loss: 0.3957 - accuracy: 0.8870 - val\_loss: 0.1669 - val\_accuracy: 0.9494

Epoch 3/20

60000/60000 [=====] - 5s 90us/step - loss: 0.2933 - accuracy: 0.9186 - val\_loss: 0.1370 - val\_accuracy: 0.9582

Epoch 4/20

60000/60000 [=====] - 6s 92us/step - loss: 0.2417 - accuracy: 0.9342 - val\_loss: 0.1212 - val\_accuracy: 0.9634

Epoch 5/20

60000/60000 [=====] - 6s 95us/step - loss: 0.2106 - accuracy: 0.9423 - val\_loss: 0.1101 - val\_accuracy: 0.9682

Epoch 6/20

60000/60000 [=====] - 6s 93us/step - loss: 0.1920 - accuracy: 0.9476 - val\_loss: 0.0988 - val\_accuracy: 0.9719

Epoch 7/20

60000/60000 [=====] - 5s 92us/step - loss: 0.1718 - accuracy: 0.9532 - val\_loss: 0.0944 - val\_accuracy: 0.9735

Epoch 8/20

60000/60000 [=====] - 6s 94us/step - loss: 0.1580 - accuracy: 0.9564 - val\_loss: 0.0860 - val\_accuracy: 0.9752

Epoch 9/20

60000/60000 [=====] - 6s 93us/step - loss: 0.1497 - accuracy: 0.9588 - val\_loss: 0.0855 - val\_accuracy: 0.9754

Epoch 10/20

60000/60000 [=====] - 6s 94us/step - loss: 0.1399 - accuracy: 0.9610 - val\_loss: 0.0832 - val\_accuracy: 0.9767

Epoch 11/20

60000/60000 [=====] - 6s 92us/step - loss: 0.1323 - accuracy: 0.9631 - val\_loss: 0.0762 - val\_accuracy: 0.9782

Epoch 12/20

60000/60000 [=====] - 6s 94us/step - loss: 0.1224 - accuracy: 0.9671 - val\_loss: 0.0734 - val\_accuracy: 0.9795

Epoch 13/20

60000/60000 [=====] - 6s 93us/step - loss: 0.1188 - accuracy: 0.9679 - val\_loss: 0.0666 - val\_accuracy: 0.9810

Epoch 14/20

60000/60000 [=====] - 6s 94us/step - loss: 0.1115 - accuracy: 0.9688 - val\_loss: 0.0703 - val\_accuracy: 0.9797

Epoch 15/20

60000/60000 [=====] - 6s 93us/step - loss: 0.1100 - accuracy: 0.9693 - val\_loss: 0.0656 - val\_accuracy: 0.9804

Epoch 16/20

60000/60000 [=====] - 6s 93us/step - loss: 0.1026 - accuracy: 0.9718 - val\_loss: 0.0705 - val\_accuracy: 0.9807

Epoch 17/20

60000/60000 [=====] - 6s 94us/step - loss: 0.1006 - accuracy: 0.9713 - val\_loss: 0.0659 - val\_accuracy: 0.9813

Epoch 18/20

```
60000/60000 [=====] - 6s 93us/step - loss: 0.0951
- accuracy: 0.9729 - val_loss: 0.0698 - val_accuracy: 0.9803
Epoch 19/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0917
- accuracy: 0.9747 - val_loss: 0.0657 - val_accuracy: 0.9808
Epoch 20/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0924
- accuracy: 0.9749 - val_loss: 0.0704 - val_accuracy: 0.9825
```



In [42]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.07038851246719714

Test accuracy: 0.9825000166893005



In [43]:

```

w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

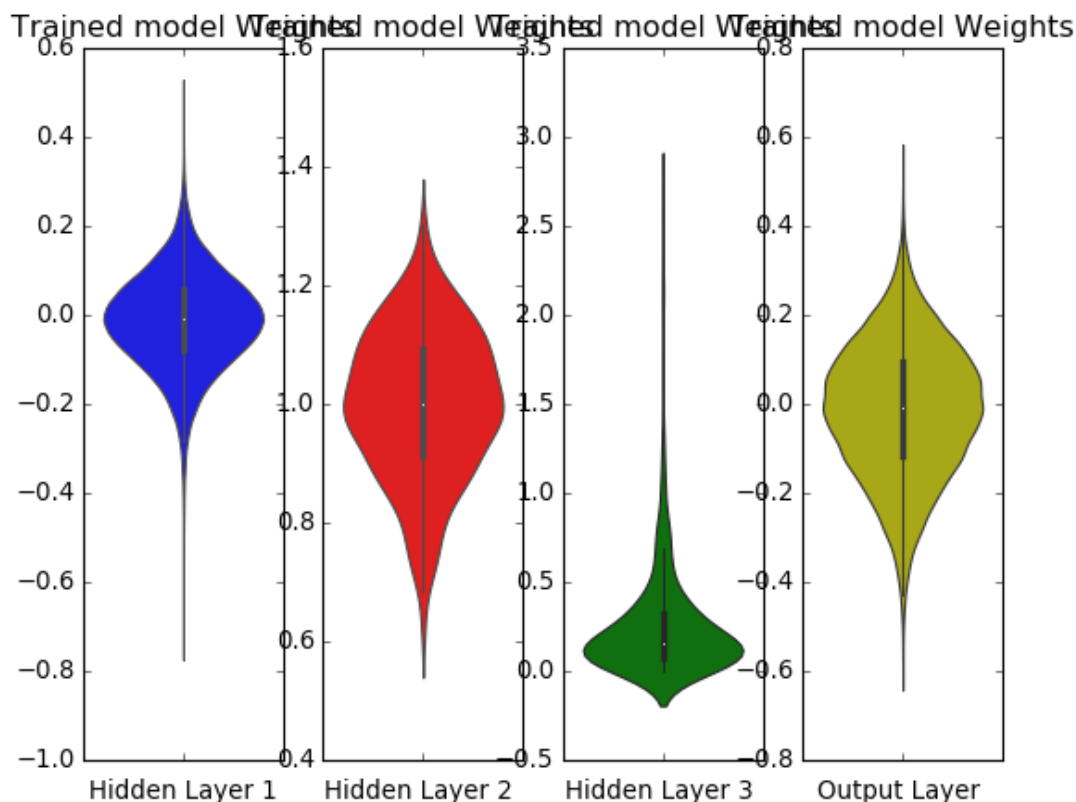
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



**MLP + ReLU + ADAM with 3 hidden layers with Dropout (dropout rate = 0.25)**

In [48]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-funct
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.25))

model_drop.add(Dense(128, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.25))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.25))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Model: "sequential\_12"

Layer (type)	Output Shape	Param #
dense_39 (Dense)	(None, 364)	285740
batch_normalization_23 (Batch Normalization)	(None, 364)	1456
dropout_18 (Dropout)	(None, 364)	0
dense_40 (Dense)	(None, 128)	46720
batch_normalization_24 (Batch Normalization)	(None, 128)	512
dropout_19 (Dropout)	(None, 128)	0
dense_41 (Dense)	(None, 52)	6708
batch_normalization_25 (Batch Normalization)	(None, 52)	208
dropout_20 (Dropout)	(None, 52)	0
dense_42 (Dense)	(None, 10)	530
Total params: 341,874		
Trainable params: 340,786		
Non-trainable params: 1,088		

In [49]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
                        validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 123us/step - loss: 0.469  
6 - accuracy: 0.8598 - val\_loss: 0.1418 - val\_accuracy: 0.9546

Epoch 2/20

60000/60000 [=====] - 6s 103us/step - loss: 0.195  
8 - accuracy: 0.9419 - val\_loss: 0.1066 - val\_accuracy: 0.9661

Epoch 3/20

60000/60000 [=====] - 6s 102us/step - loss: 0.146  
5 - accuracy: 0.9567 - val\_loss: 0.0860 - val\_accuracy: 0.9710

Epoch 4/20

60000/60000 [=====] - 6s 101us/step - loss: 0.117  
5 - accuracy: 0.9639 - val\_loss: 0.0793 - val\_accuracy: 0.9748

Epoch 5/20

60000/60000 [=====] - 6s 101us/step - loss: 0.098  
9 - accuracy: 0.9702 - val\_loss: 0.0707 - val\_accuracy: 0.9779

Epoch 6/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0889  
- accuracy: 0.9737 - val\_loss: 0.0717 - val\_accuracy: 0.9786

Epoch 7/20

60000/60000 [=====] - 6s 102us/step - loss: 0.079  
1 - accuracy: 0.9751 - val\_loss: 0.0709 - val\_accuracy: 0.9774

Epoch 8/20

60000/60000 [=====] - 6s 103us/step - loss: 0.070  
6 - accuracy: 0.9783 - val\_loss: 0.0690 - val\_accuracy: 0.9801

Epoch 9/20

60000/60000 [=====] - 6s 100us/step - loss: 0.062  
8 - accuracy: 0.9807 - val\_loss: 0.0716 - val\_accuracy: 0.9802

Epoch 10/20

60000/60000 [=====] - 6s 97us/step - loss: 0.0590  
- accuracy: 0.9816 - val\_loss: 0.0696 - val\_accuracy: 0.9800

Epoch 11/20

60000/60000 [=====] - 6s 100us/step - loss: 0.058  
1 - accuracy: 0.9820 - val\_loss: 0.0673 - val\_accuracy: 0.9813

Epoch 12/20

60000/60000 [=====] - 6s 100us/step - loss: 0.048  
8 - accuracy: 0.9846 - val\_loss: 0.0659 - val\_accuracy: 0.9811

Epoch 13/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0477  
- accuracy: 0.9854 - val\_loss: 0.0665 - val\_accuracy: 0.9817

Epoch 14/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0460  
- accuracy: 0.9852 - val\_loss: 0.0636 - val\_accuracy: 0.9830

Epoch 15/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0448  
- accuracy: 0.9861 - val\_loss: 0.0642 - val\_accuracy: 0.9833

Epoch 16/20

60000/60000 [=====] - 6s 100us/step - loss: 0.042  
9 - accuracy: 0.9863 - val\_loss: 0.0648 - val\_accuracy: 0.9816

Epoch 17/20

60000/60000 [=====] - 6s 100us/step - loss: 0.038  
9 - accuracy: 0.9871 - val\_loss: 0.0623 - val\_accuracy: 0.9831

Epoch 18/20

```
60000/60000 [=====] - 6s 99us/step - loss: 0.0358  
- accuracy: 0.9888 - val_loss: 0.0660 - val_accuracy: 0.9828  
Epoch 19/20  
60000/60000 [=====] - 6s 98us/step - loss: 0.0381  
- accuracy: 0.9876 - val_loss: 0.0652 - val_accuracy: 0.9833  
Epoch 20/20  
60000/60000 [=====] - 6s 99us/step - loss: 0.0362  
- accuracy: 0.9885 - val_loss: 0.0638 - val_accuracy: 0.9820
```

In [50]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06377539714767481  
 Test accuracy: 0.9819999933242798



In [51]:

```

w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```





## **MLP + ReLU + ADAM with 5 hidden layers without Dropout and Batch Normalisation**

In [52]:

```
model_relu = Sequential()
model_relu.add(Dense(364, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_relu.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_relu.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_relu.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)
```

Model: "sequential\_13"

Layer (type)	Output Shape	Param #
dense_43 (Dense)	(None, 364)	285740
dense_44 (Dense)	(None, 128)	46720
dense_45 (Dense)	(None, 64)	8256
dense_46 (Dense)	(None, 32)	2080
dense_47 (Dense)	(None, 16)	528
dense_48 (Dense)	(None, 10)	170
Total params: 343,494		
Trainable params: 343,494		
Non-trainable params: 0		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 79us/step - loss: 0.3766 - accuracy: 0.8903 - val\_loss: 0.1837 - val\_accuracy: 0.9487

Epoch 2/20

60000/60000 [=====] - 4s 67us/step - loss: 0.1257 - accuracy: 0.9641 - val\_loss: 0.1208 - val\_accuracy: 0.9634

Epoch 3/20

60000/60000 [=====] - 4s 67us/step - loss: 0.0797 - accuracy: 0.9764 - val\_loss: 0.1173 - val\_accuracy: 0.9648

Epoch 4/20

60000/60000 [=====] - 4s 67us/step - loss: 0.0601 - accuracy: 0.9812 - val\_loss: 0.1019 - val\_accuracy: 0.9697

Epoch 5/20

60000/60000 [=====] - 4s 67us/step - loss: 0.0433 - accuracy: 0.9869 - val\_loss: 0.0933 - val\_accuracy: 0.9751

Epoch 6/20

60000/60000 [=====] - 4s 64us/step - loss: 0.0362 - accuracy: 0.9883 - val\_loss: 0.0842 - val\_accuracy: 0.9771

Epoch 7/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0316 - accuracy: 0.9899 - val\_loss: 0.0870 - val\_accuracy: 0.9763

Epoch 8/20

```
60000/60000 [=====] - 4s 65us/step - loss: 0.0239 -  
accuracy: 0.9920 - val_loss: 0.0919 - val_accuracy: 0.9753  
Epoch 9/20  
60000/60000 [=====] - 4s 66us/step - loss: 0.0269 -  
accuracy: 0.9913 - val_loss: 0.1013 - val_accuracy: 0.9760  
Epoch 10/20  
60000/60000 [=====] - 4s 64us/step - loss: 0.0257 -  
accuracy: 0.9910 - val_loss: 0.1084 - val_accuracy: 0.9755  
Epoch 11/20  
60000/60000 [=====] - 4s 65us/step - loss: 0.0176 -  
accuracy: 0.9947 - val_loss: 0.1006 - val_accuracy: 0.9777  
Epoch 12/20  
60000/60000 [=====] - 4s 65us/step - loss: 0.0163 -  
accuracy: 0.9947 - val_loss: 0.1221 - val_accuracy: 0.9752  
Epoch 13/20  
60000/60000 [=====] - 4s 66us/step - loss: 0.0210 -  
accuracy: 0.9934 - val_loss: 0.1091 - val_accuracy: 0.9754  
Epoch 14/20  
60000/60000 [=====] - 4s 65us/step - loss: 0.0168 -  
accuracy: 0.9944 - val_loss: 0.1114 - val_accuracy: 0.9762  
Epoch 15/20  
60000/60000 [=====] - 4s 66us/step - loss: 0.0166 -  
accuracy: 0.9946 - val_loss: 0.1143 - val_accuracy: 0.9749  
Epoch 16/20  
60000/60000 [=====] - 4s 67us/step - loss: 0.0116 -  
accuracy: 0.9962 - val_loss: 0.1110 - val_accuracy: 0.9798  
Epoch 17/20  
60000/60000 [=====] - 4s 64us/step - loss: 0.0133 -  
accuracy: 0.9959 - val_loss: 0.0991 - val_accuracy: 0.9784  
Epoch 18/20  
60000/60000 [=====] - 4s 63us/step - loss: 0.0166 -  
accuracy: 0.9947 - val_loss: 0.0953 - val_accuracy: 0.9801  
Epoch 19/20  
60000/60000 [=====] - 4s 66us/step - loss: 0.0103 -  
accuracy: 0.9968 - val_loss: 0.1048 - val_accuracy: 0.9776  
Epoch 20/20  
60000/60000 [=====] - 4s 66us/step - loss: 0.0112 -  
accuracy: 0.9966 - val_loss: 0.1061 - val_accuracy: 0.9777
```

In [53]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.1060613916828221

Test accuracy: 0.9776999950408936





In [54]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='cyan')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='brown')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



**MLP + ReLU + ADAM with 5 hidden layers with Batch Normalisation**

In [55]:

```

# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(0, \sigma) = N(0, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(0, \sigma) = N(0, 0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(364, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.120, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.055, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```

Model: "sequential\_14"

Layer (type)	Output Shape	Param #
dense_49 (Dense)	(None, 364)	285740
batch_normalization_26 (Batch Normalization)	(None, 364)	1456
dense_50 (Dense)	(None, 128)	46720
batch_normalization_27 (Batch Normalization)	(None, 128)	512
dense_51 (Dense)	(None, 64)	8256
batch_normalization_28 (Batch Normalization)	(None, 64)	256
dense_52 (Dense)	(None, 32)	2080
batch_normalization_29 (Batch Normalization)	(None, 32)	128
dense_53 (Dense)	(None, 16)	528
batch_normalization_30 (Batch Normalization)	(None, 16)	64
dense_54 (Dense)	(None, 10)	170



```
=====
Total params: 345,910
Trainable params: 344,702
Non-trainable params: 1,208

```

---

In [56]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                          validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 133us/step - loss: 0.4110  
- accuracy: 0.8947 - val\_loss: 0.1524 - val\_accuracy: 0.9563

Epoch 2/20

60000/60000 [=====] - 6s 101us/step - loss: 0.1170  
- accuracy: 0.9677 - val\_loss: 0.1063 - val\_accuracy: 0.9708

Epoch 3/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0744 -  
accuracy: 0.9780 - val\_loss: 0.0919 - val\_accuracy: 0.9743

Epoch 4/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0561 -  
accuracy: 0.9832 - val\_loss: 0.0901 - val\_accuracy: 0.9741

Epoch 5/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0473 -  
accuracy: 0.9853 - val\_loss: 0.1002 - val\_accuracy: 0.9725

Epoch 6/20

60000/60000 [=====] - 6s 101us/step - loss: 0.0354 -  
accuracy: 0.9887 - val\_loss: 0.0868 - val\_accuracy: 0.9762

Epoch 7/20

60000/60000 [=====] - 6s 97us/step - loss: 0.0310 -  
accuracy: 0.9905 - val\_loss: 0.0941 - val\_accuracy: 0.9741

Epoch 8/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0266 -  
accuracy: 0.9915 - val\_loss: 0.0865 - val\_accuracy: 0.9769

Epoch 9/20

60000/60000 [=====] - 6s 97us/step - loss: 0.0252 -  
accuracy: 0.9916 - val\_loss: 0.0849 - val\_accuracy: 0.9761

Epoch 10/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0258 -  
accuracy: 0.9917 - val\_loss: 0.0921 - val\_accuracy: 0.9756

Epoch 11/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0204 -  
accuracy: 0.9933 - val\_loss: 0.0846 - val\_accuracy: 0.9771

Epoch 12/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0202 -  
accuracy: 0.9936 - val\_loss: 0.0828 - val\_accuracy: 0.9783

Epoch 13/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0147 -  
accuracy: 0.9951 - val\_loss: 0.0908 - val\_accuracy: 0.9764

Epoch 14/20

60000/60000 [=====] - 6s 97us/step - loss: 0.0161 -  
accuracy: 0.9946 - val\_loss: 0.0777 - val\_accuracy: 0.9807

Epoch 15/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0148 -  
accuracy: 0.9949 - val\_loss: 0.1032 - val\_accuracy: 0.9747

Epoch 16/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0137 -  
accuracy: 0.9956 - val\_loss: 0.0886 - val\_accuracy: 0.9780

Epoch 17/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0128 -  
accuracy: 0.9959 - val\_loss: 0.0939 - val\_accuracy: 0.9780

Epoch 18/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0123

- accuracy: 0.9959 - val\_loss: 0.1088 - val\_accuracy: 0.9741

Epoch 19/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0130 -

accuracy: 0.9956 - val\_loss: 0.0745 - val\_accuracy: 0.9827

Epoch 20/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0093

- accuracy: 0.9970 - val\_loss: 0.0868 - val\_accuracy: 0.9798

In [57]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.1060613916828221

Test accuracy: 0.9776999950408936



In [58]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='cyan')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='brown')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



**MLP + ReLU + ADAM with 5 hidden layers with Dropout (dropout rate = 0.5)**

In [59]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-funct
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Model: "sequential\_15"

Layer (type)	Output Shape	Param #
dense_55 (Dense)	(None, 364)	285740
dropout_21 (Dropout)	(None, 364)	0
dense_56 (Dense)	(None, 128)	46720
dropout_22 (Dropout)	(None, 128)	0
dense_57 (Dense)	(None, 64)	8256
dropout_23 (Dropout)	(None, 64)	0
dense_58 (Dense)	(None, 32)	2080
dropout_24 (Dropout)	(None, 32)	0
dense_59 (Dense)	(None, 16)	528
dropout_25 (Dropout)	(None, 16)	0
dense_60 (Dense)	(None, 10)	170
Total params: 343,494		
Trainable params: 343,494		
Non-trainable params: 0		

In [60]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
                        validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 98us/step - loss: 2.7044 - accuracy: 0.1195 - val\_loss: 2.3014 - val\_accuracy: 0.1136

Epoch 2/20

60000/60000 [=====] - 5s 83us/step - loss: 2.1709 - accuracy: 0.2078 - val\_loss: 1.8698 - val\_accuracy: 0.3667

Epoch 3/20

60000/60000 [=====] - 5s 80us/step - loss: 1.8653 - accuracy: 0.3251 - val\_loss: 1.5882 - val\_accuracy: 0.4043

Epoch 4/20

60000/60000 [=====] - 5s 80us/step - loss: 1.6584 - accuracy: 0.4065 - val\_loss: 1.4147 - val\_accuracy: 0.5350

Epoch 5/20

60000/60000 [=====] - 5s 80us/step - loss: 1.4961 - accuracy: 0.4598 - val\_loss: 1.2402 - val\_accuracy: 0.5695

Epoch 6/20

60000/60000 [=====] - 5s 81us/step - loss: 1.3807 - accuracy: 0.4925 - val\_loss: 1.1266 - val\_accuracy: 0.6056

Epoch 7/20

60000/60000 [=====] - 5s 83us/step - loss: 1.2763 - accuracy: 0.5304 - val\_loss: 0.9972 - val\_accuracy: 0.6800

Epoch 8/20

60000/60000 [=====] - 5s 81us/step - loss: 1.1909 - accuracy: 0.5656 - val\_loss: 0.8996 - val\_accuracy: 0.7056

Epoch 9/20

60000/60000 [=====] - 5s 81us/step - loss: 1.1053 - accuracy: 0.6075 - val\_loss: 0.8130 - val\_accuracy: 0.7336

Epoch 10/20

60000/60000 [=====] - 6s 93us/step - loss: 1.0238 - accuracy: 0.6384 - val\_loss: 0.7449 - val\_accuracy: 0.7680

Epoch 11/20

60000/60000 [=====] - 6s 92us/step - loss: 0.9651 - accuracy: 0.6621 - val\_loss: 0.6997 - val\_accuracy: 0.8079

Epoch 12/20

60000/60000 [=====] - 5s 83us/step - loss: 0.9139 - accuracy: 0.6915 - val\_loss: 0.5961 - val\_accuracy: 0.8414

Epoch 13/20

60000/60000 [=====] - 5s 82us/step - loss: 0.8529 - accuracy: 0.7160 - val\_loss: 0.5478 - val\_accuracy: 0.8608

Epoch 14/20

60000/60000 [=====] - 5s 81us/step - loss: 0.8207 - accuracy: 0.7338 - val\_loss: 0.5052 - val\_accuracy: 0.8943

Epoch 15/20

60000/60000 [=====] - 5s 81us/step - loss: 0.7725 - accuracy: 0.7565 - val\_loss: 0.4904 - val\_accuracy: 0.8933

Epoch 16/20

60000/60000 [=====] - 5s 83us/step - loss: 0.7675 - accuracy: 0.7672 - val\_loss: 0.4819 - val\_accuracy: 0.9102

Epoch 17/20

60000/60000 [=====] - 5s 81us/step - loss: 0.7231 - accuracy: 0.7794 - val\_loss: 0.4629 - val\_accuracy: 0.9135

Epoch 18/20

60000/60000 [=====] - 5s 79us/step - loss: 0.6916 -



accuracy: 0.7927 - val\_loss: 0.4751 - val\_accuracy: 0.9215

Epoch 19/20

60000/60000 [=====] - 5s 79us/step - loss: 0.6795 -

accuracy: 0.7985 - val\_loss: 0.5045 - val\_accuracy: 0.9202

Epoch 20/20

60000/60000 [=====] - 5s 84us/step - loss: 0.6536 -

accuracy: 0.8087 - val\_loss: 0.4650 - val\_accuracy: 0.9332

In [61]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.1060613916828221

Test accuracy: 0.9776999950408936



In [62]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='cyan')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='brown')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



**MLP + ReLU + ADAM with 5 hidden layers with Dropout (dropout rate = 0.25)**

In [63]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-funct
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(Dropout(0.25))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(Dropout(0.25))

model_drop.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(Dropout(0.25))

model_drop.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(Dropout(0.25))

model_drop.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_drop.add(Dropout(0.25))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Model: "sequential\_16"

Layer (type)	Output Shape	Param #
dense_61 (Dense)	(None, 364)	285740
dropout_26 (Dropout)	(None, 364)	0
dense_62 (Dense)	(None, 128)	46720
dropout_27 (Dropout)	(None, 128)	0
dense_63 (Dense)	(None, 64)	8256
dropout_28 (Dropout)	(None, 64)	0
dense_64 (Dense)	(None, 32)	2080
dropout_29 (Dropout)	(None, 32)	0
dense_65 (Dense)	(None, 16)	528
dropout_30 (Dropout)	(None, 16)	0
dense_66 (Dense)	(None, 10)	170
Total params: 343,494		
Trainable params: 343,494		
Non-trainable params: 0		

In [64]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
                        validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 105us/step - loss: 1.4585  
- accuracy: 0.5030 - val\_loss: 0.3698 - val\_accuracy: 0.9152

Epoch 2/20

60000/60000 [=====] - 5s 82us/step - loss: 0.5783 -  
accuracy: 0.8296 - val\_loss: 0.2133 - val\_accuracy: 0.9456

Epoch 3/20

60000/60000 [=====] - 5s 83us/step - loss: 0.4104 -  
accuracy: 0.8878 - val\_loss: 0.1751 - val\_accuracy: 0.9563

Epoch 4/20

60000/60000 [=====] - 5s 83us/step - loss: 0.3233 -  
accuracy: 0.9148 - val\_loss: 0.1400 - val\_accuracy: 0.9671

Epoch 5/20

60000/60000 [=====] - 5s 84us/step - loss: 0.2796 -  
accuracy: 0.9281 - val\_loss: 0.1321 - val\_accuracy: 0.9671

Epoch 6/20

60000/60000 [=====] - 5s 82us/step - loss: 0.2445 -  
accuracy: 0.9377 - val\_loss: 0.1359 - val\_accuracy: 0.9683

Epoch 7/20

60000/60000 [=====] - 5s 81us/step - loss: 0.2239 -  
accuracy: 0.9448 - val\_loss: 0.1203 - val\_accuracy: 0.9719

Epoch 8/20

60000/60000 [=====] - 5s 82us/step - loss: 0.1987 -  
accuracy: 0.9500 - val\_loss: 0.1399 - val\_accuracy: 0.9725

Epoch 9/20

60000/60000 [=====] - 5s 83us/step - loss: 0.1833 -  
accuracy: 0.9551 - val\_loss: 0.1320 - val\_accuracy: 0.9730

Epoch 10/20

60000/60000 [=====] - 5s 80us/step - loss: 0.1727 -  
accuracy: 0.9566 - val\_loss: 0.1266 - val\_accuracy: 0.9747

Epoch 11/20

60000/60000 [=====] - 5s 82us/step - loss: 0.1652 -  
accuracy: 0.9588 - val\_loss: 0.1348 - val\_accuracy: 0.9733

Epoch 12/20

60000/60000 [=====] - 5s 85us/step - loss: 0.1513 -  
accuracy: 0.9623 - val\_loss: 0.1096 - val\_accuracy: 0.9768

Epoch 13/20

60000/60000 [=====] - 5s 82us/step - loss: 0.1451 -  
accuracy: 0.9639 - val\_loss: 0.1105 - val\_accuracy: 0.9778

Epoch 14/20

60000/60000 [=====] - 5s 80us/step - loss: 0.1326 -  
accuracy: 0.9662 - val\_loss: 0.1034 - val\_accuracy: 0.9805

Epoch 15/20

60000/60000 [=====] - 5s 80us/step - loss: 0.1343 -  
accuracy: 0.9657 - val\_loss: 0.1047 - val\_accuracy: 0.9790

Epoch 16/20

60000/60000 [=====] - 5s 83us/step - loss: 0.1284 -  
accuracy: 0.9679 - val\_loss: 0.1215 - val\_accuracy: 0.9792

Epoch 17/20

60000/60000 [=====] - 5s 80us/step - loss: 0.1253 -  
accuracy: 0.9683 - val\_loss: 0.1252 - val\_accuracy: 0.9801

Epoch 18/20

60000/60000 [=====] - 5s 80us/step - loss: 0.1126 -

accuracy: 0.9715 - val\_loss: 0.1112 - val\_accuracy: 0.9800

Epoch 19/20

60000/60000 [=====] - 5s 79us/step - loss: 0.1124 -

accuracy: 0.9712 - val\_loss: 0.1283 - val\_accuracy: 0.9788

Epoch 20/20

60000/60000 [=====] - 5s 80us/step - loss: 0.1062 -

accuracy: 0.9738 - val\_loss: 0.1156 - val\_accuracy: 0.9803

In [65]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.1060613916828221

Test accuracy: 0.9776999950408936

/home/komalumrethe/anaconda3/lib/python3.5/site-packages/matplotlib/pyplot.py:524: RuntimeWarning: More than 20 figures have been opened. Figures created through the pyplot interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max\_open\_warning`).  
max\_open\_warning, RuntimeWarning)





In [66]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='cyan')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='brown')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

/home/komalumrethe/anaconda3/lib/python3.5/site-packages/matplotlib/pyplot.p  
y:524: RuntimeWarning: More than 20 figures have been opened. Figures create  
d through the pyplot interface (`matplotlib.pyplot.figure`) are retained unt  
il explicitly closed and may consume too much memory. (To control this warni  
ng, see the rcParam `figure.max\_open\_warning`).  
max\_open\_warning, RuntimeWarning)



In [67]:

```
from prettytable import PrettyTable

x = PrettyTable()
x.field_names = ["Architecture", "parameters", "Accuracy"]

x.add_row(["2 layer", "Without Batch Normalization and Dropout", 97.92])
x.add_row(["2 layer", "With Batch Normalization", 97.96])
x.add_row(["2 layer", "With Dropuot(dropout rate = 0.5)", 98.28])
x.add_row(["2 layer", "With Dropuot(dropout rate = 0.25)", 98.29])
x.add_row(["2 layer", "With Dropuot(dropout rate = 0.75)", 98.35])

x.add_row(["3 layer", "Without Batch Normalization and Dropout", 97.90])
x.add_row(["3 layer", "With Batch Normalization", 97.89])
x.add_row(["3 layer", "With Dropuot(dropout rate = 0.5)", 98.25])
x.add_row(["3 layer", "With Dropuot(dropout rate = 0.25)", 98.19])

x.add_row(["5 layer", "Without Batch Normalization and Dropout", 97.7699])
x.add_row(["5 layer", "With Batch Normalization", 97.7699])
x.add_row(["5 layer", "With Dropuot(dropout rate = 0.5)", 97.7699])
x.add_row(["5 layer", "With Dropuot(dropout rate = 0.25)", 97.7699])

print(x)
```

Architecture	parameters	Accuracy
2 layer	Without Batch Normalization and Dropout	97.92
2 layer	With Batch Normalization	97.96
2 layer	With Dropuot(dropout rate = 0.5)	98.28
2 layer	With Dropuot(dropout rate = 0.25)	98.29
2 layer	With Dropuot(dropout rate = 0.75)	98.35
3 layer	Without Batch Normalization and Dropout	97.9
3 layer	With Batch Normalization	97.89
3 layer	With Dropuot(dropout rate = 0.5)	98.25
3 layer	With Dropuot(dropout rate = 0.25)	98.19
5 layer	Without Batch Normalization and Dropout	97.7699
5 layer	With Batch Normalization	97.7699
5 layer	With Dropuot(dropout rate = 0.5)	97.7699
5 layer	With Dropuot(dropout rate = 0.25)	97.7699

## Procedure Followed

- Flattened the 28\*28 dimensional MNIST data to 784
- Normalized the data
- Used a softmax classifier of output dimensions = 10
- Created multiple models in Keras with various parameter combinations like activation function = 'relu', optimizer = 'Adam', with/without dropout of different rates and with/without Batch normalization
- Plotted the epoch vs Train/Test loss of each model
- Plotted the weights of each model

